

OpenCL exercise 1: Basics

Steffen Kieß

Task 1

- ▶ Calculate cosine function for a number of values on the GPU
- ▶ CPU code:

```
for (std::size_t i = 0; i < h_output.size (); i++)  
    h_output[i] = std::cos (h_input[i]);
```

- ▶ Use one work item per value on the GPU

Steps

- ▶ Allocate memory on the device (`cl::Buffer` constructor)
- ▶ Initialize the memory on the device
`(cl::CommandQueue::enqueueWriteBuffer())`
- ▶ Copy the input data to the device
`(cl::CommandQueue::enqueueWriteBuffer())`
- ▶ Launch the kernel (`cl::Kernel::setArg()`,
`cl::CommandQueue::enqueueNDRangeKernel()`)
- ▶ Copy the output data to the host
`(cl::CommandQueue::enqueueReadBuffer())`

Task 2

Add code for measuring

- ▶ Time needed for calculation on host
- ▶ Time needed for calculation on device
- ▶ Time needed for memory transactions

Speedup:

- ▶ $\text{Time on the CPU} / \text{Time on the GPU}$
- ▶ $\text{Time on the CPU} / (\text{Time on the GPU} + \text{Time for Memory transactions})$

Compare the times using a Debug build and a Release build.

Task 3

Use `native_cos()` instead of `cos()` in the kernel and compare the performance.

Syntax: Memory allocation

Allocate memory:

```
cl::Buffer::Buffer(cl::Context context,  
    cl_mem_flags flags, std::size_t size);
```

context = The context to use

flags = Normally CL_MEM_READ_WRITE

size = Number of bytes (not elements) to allocate

Syntax: Copying data

Copy data from CPU to GPU (global) memory:

```
cl::CommandQueue::enqueueWriteBuffer(cl::Buffer buffer,  
    bool blocking, std::size_t offset, std::size_t size,  
    const void* ptr, eventsToWaitFor = NULL,  
    cl::Event* resultEvent = NULL) const;
```

buffer = The buffer to copy to

blocking = Wait until the copy operation has finished (normally true)

offset = Offset into the buffer (in bytes)

size = Number of bytes (not elements) to copy

ptr = Pointer to source data in CPU memory

eventsToWaitFor = Events which have to occur before the copy operation is started, normally NULL

resultEvent = Pointer to a variable where an event is stored (can be used for profiling)

Syntax: Copying data

Copy data from GPU (global) to CPU memory:

```
cl::CommandQueue::enqueueReadBuffer(cl::Buffer buffer,  
    bool blocking, std::size_t offset, std::size_t size,  
    void* ptr, eventsToWaitFor = NULL,  
    cl::Event* resultEvent = NULL) const;
```

buffer = The buffer to copy from

blocking = Wait until the copy operation has finished (normally true)

offset = Offset into the buffer (in bytes)

size = Number of bytes (not elements) to copy

ptr = Pointer to destination in CPU memory

eventsToWaitFor = Events which have to occur before the copy operation is started, normally NULL

resultEvent = Pointer to a variable where an event is stored (can be used for profiling)

Syntax: Launching a kernel

Set parameters for a kernel launch:

```
cl::Kernel::setArg<T>(cl_uint index, T value);
```

T = The type of the parameter (e.g. `cl_int` or `cl::Buffer`)

index = 0-based index of the parameter

value = The value to use for the parameter

Launch the kernel:

```
cl::CommandQueue::enqueueNDRangeKernel(Kernel kernel,  
    NDRange offset, NDRange global, NDRange local,  
    eventsToWaitFor = NULL, cl::Event* event=NULL) const;
```

kernel = The kernel to launch

offset = Normally 0 or `cl::NullRange`

global = The overall number of work items

local = The number of work items per work group

Syntax: Kernel code

Declaring pointers:

```
__global int* foo; // Declare foo as a pointer to global mem.  
__local int* foo; // ... to local memory  
__private int* foo; // ... to private memory  
__constant int* foo; // ... to constant memory
```

Get global index of the current work item in the x-direction:

```
size_t i = get_global_id(0);
```

Profiling

On the CPU:

```
Core::TimeSpan time1 = Core::getCurrentTime();  
// Execute some code ...  
Core::TimeSpan time2 = Core::getCurrentTime();  
  
Core::TimeSpan time = time2 - time1;  
std::cout << time << std::endl;
```

On the GPU:

```
cl::Event event;  
queue.enqueue(..., &event);  
queue.finish(); // or enqueue*Buffer() with blocking = true  
  
Core::TimeSpan time = OpenCL::getElapsedTime(event);  
std::cout << time << std::endl;
```